

Lab 1 – Introduction to MPI

Issues with shared memory programming

- Parallel tasks are run by threads.
- All threads live on the same node & share memory.
- Bottleneck: Resources on the single node.
- Overhead due to creation & deletion of threads.
- Can lead to race conditions.

Distributed Memory Programming

- Parallel tasks are processes.
- Each process has its own private memory.
- Processes need not be on the same node.
- Possibility of introducing new bugs like deadlocks.
- **MUST EXPLICITLY CODE IN THE COMMUNICATION BETWEEN PROCESSES: MPI (MESSAGE PASSING INTERFACE).**

What is MPI?

- MPI defines a standard library for message-passing that can be used to develop portable message-passing programs
- Uses either C or Fortran
- Defines the **syntax** and **semantics** of a core set of library routines
- Allows data to be passed between processes in a distributed memory environment
- Distributed memory
 - Each processor has local memory
 - Cannot directly access the memory of other processors

MPI Forum

- First message-passing interface standard
 - Successor to PVM (Parallel Virtual Machine)
- Sixty people from forty different organizations (academia and industry)
- International representation
- MPI 1.1 Standard developed from 92-94
- MPI 2.0 Standard developed from 95-97
- Standards documents
 - <http://www.mcs.anl.gov/mpi/index.html>
 - <http://www.mpi-forum.org/docs/docs.html> (postscript versions)

Goals and Scope

- MPI's prime goals are:
 - To provide source-code portability
 - To allow efficient implementation
- It also offers:
 - A great deal of functionality
 - Support for heterogeneous parallel architectures

MPI Library

- Over 125 routines
- However, able to write fully-functional parallel programs using only 6 of them:
 - MPI_Init Initializes MPI
 - MPI_Finalize Terminates MPI
 - MPI_Comm_size Determines no of processes
 - MPI_Comm_rank Determines label of calling process
 - MPI_Send Sends a message
 - MPI_Recv Receives a message

Starting the MPI Library

- **MPI_Init**

- Called only once prior to any calls to other MPI routines
- Initializes the MPI environment
- Must be called by all processes

int MPI_Init(int *argc, char *argv)**

- argc and argv are the command line arguments of the C program
- Upon successful execution, MPI_Init returns MPI_SUCCESS

Terminating the MPI Library

- **MPI_Finalize()**

- Called at the end of the computation
- Clean-up tasks to terminate the MPI environment
- No MPI calls after MPI_Finalize
- Must be called by all processes

int MPI_Finalize()

- Upon successful execution, returns MPI_SUCCESS

Naming Practices & Argument Conventions

- All MPI routines, data types and constants are prefixed by “**MPI_**”
- MPI constants and data structures are defined for C in the file “mpi.h”.
- This header file must be included in each MPI program

Communicators

- Communication domain: set of processes that are allowed to communicate with each other
- Information about communication domains are stored in variables of type `MPI_Comm`, called **communicators**
- Used as arguments to all message transfer MPI routines
- Each process can belong to many different communication domains

MPI_COMM_WORLD

- In general, all processes may need to communicate with each other
- Default communicator, MPI_COMM_WORLD, which includes all processes involved in the parallel execution
- However, to perform communication only within a particular group of processes, use a communicator for each such group: no messages will interfere with messages destined to other groups (Using MPI_Group)

Find number of processes

- `MPI_Comm_size`

`int MPI_Comm_size(MPI_Comm comm, int *size)`

- Returns in the variable `size` the number of processes that belong to the communicator `comm`

e.g.

`int size;`

`MPI_Comm_size(MPI_COMM_WORLD, &size);`

Rank identification

- MPI_Comm_rank

int MPI_Comm_rank(MPI_Comm comm, int *rank)

- Every process is uniquely identified by its rank
- Starts with zero and goes to (n-1) where n is the number of processes requested
- On return, the variable rank stores the rank of the process

e.g.

int rank;

MPI_Comm_rank(MPI_COMM_WORLD, &rank)

Hello World

```
#include <mpi.h>

main(int argc, char *argv[])
{
    int npes, myrank;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &npes);
    MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
    printf("From process %d out of %d, Hello
        World!\n", myrank, npes);
    MPI_Finalize();
}
```